

# aishba-waqar-46997-lab-11

April 22, 2025

```
[23]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.naive_bayes import GaussianNB
      from sklearn.metrics import accuracy_score, classification_report
      from sklearn.preprocessing import LabelEncoder, OneHotEncoder

[24]: df = pd.read_csv('/content/Public Livelihood Data.csv')
      df.dropna(inplace=True)

[25]: df.columns

[25]: Index(['Designation', 'Education', 'Marital Status', 'Field', 'Race', 'Gender',
          'Country', 'Salary'],
          dtype='object')

[34]: # Create a LabelEncoder object
      le = LabelEncoder()

      # Apply LabelEncoder to only categorical columns
      categorical_columns = df.select_dtypes(include=['object']).columns
      df[categorical_columns] = df[categorical_columns].apply(le.fit_transform)

[35]: X = df.drop('Salary', axis=1)
      y = df['Salary']

[36]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      random_state=42)

[37]: # Step 6: Train the Naive Bayes Classifier
      model = GaussianNB()
      model.fit(X_train, y_train)

[37]: GaussianNB()

[38]: y_pred = model.predict(X_test)

[39]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Designation            32561 non-null  int64
1   Education              32561 non-null  int64
2   Marital Status         32561 non-null  int64
3   Field                  32561 non-null  int64
4   Race                   32561 non-null  int64
5   Gender                 32561 non-null  int64
6   Country                32561 non-null  int64
7   Salary                 32561 non-null  int64
dtypes: int64(8)
memory usage: 2.0 MB

```

```

[40]: accuracy = accuracy_score(y_test, y_pred)
      print("Accuracy:", accuracy)

```

Accuracy: 0.760325502840473

```

[41]: print("\nPredicted labels:", y_pred)
      print("Actual labels:   ", y_test.values)
      print("\nClassification Report:\n", classification_report(y_test,y_pred))

```

```

Predicted labels: [0 0 0 ... 1 0 0]
Actual labels:    [0 0 1 ... 1 0 0]

```

```

Classification Report:

```

	precision	recall	f1-score	support
0	0.87	0.80	0.84	4942
1	0.50	0.63	0.56	1571
accuracy			0.76	6513
macro avg	0.69	0.72	0.70	6513
weighted avg	0.78	0.76	0.77	6513